

Structuring Unstructured Data with LLMs: DSPy Practical Assignment

Why This Matters

Unstructured data (text, PDFs, web content) makes up **80-90% of data**, yet it's unusable for analysis without costly manual processing. **LLMs solve this by acting as "structured data compilers"** – converting messy text into clean, query able formats (entities, relations, graphs). This assignment tests your ability to:

- **Bridge theory and practice** (real-world data ≠ textbook examples)
- **Handle LLM uncertainty** (confidence loops, error resilience)
- **Build production-ready pipelines** (not just one-off scripts)

The provided DSPy codebase demonstrates **exactly what modern data engineering teams need**:

- 1. **Entity extraction** → Turning text into typed objects (e.g., "pelletized frass" → Drug)
- 2. **Intelligent deduplication** → Solving real-world noise (e.g., "PB IC", "pea-barley intercrop", "pea-barley intercrops" → 1 entity)
- 3. **Knowledge graph generation** → Creating visual, query able relationships (Mermaid diagrams)

Code Walkthrough (Key Concepts for Applicants)

Code Sample: https://colab.research.google.com/drive/1bw8EBRn_bRCysFnnBNUXCqg05KR8VsX?usp=sharing

1. Entity Extraction (ExtractEntities Signature)

```
class EntityWithAttr(BaseModel):
    entity: str = Field(description="the named entity")
    attr_type: str = Field(description="semantic type (e.g. Drug, Disease)")
```

```
class ExtractEntities(dspy.Signature):  
    paragraph: str = dspy.InputField()  
    entities: List[EntityWithAttr] = dspy.OutputField()
```

2 / 4

- **Why it's clever:** Uses Pydantic to **force structured outputs** from LLMs. No more regex parsing of free-text responses!
- **Your takeaway:** Always define *exactly* what the LLM should output. DSPy validates responses against your schema.

2. Deduplication with Confidence Loops

```
def deduplicate_with_lm(items, batch_size=10, target_confidence=0.9):  
    while True:  
        pred = dedup_predictor(items=batch)  
        if pred.confidence >= target_confidence: # Critical safety check!  
            return pred.deduplicated
```

- **Why it's clever:** LLMs hallucinate. This loop **self-corrects** until confidence $\geq 90\%$.
- **Your takeaway:** Never trust a single LLM call. *Always* add validation loops for critical tasks.

3. Mermaid Graph Generation

```
def triples_to_mermaid(triples, entity_list):  
    # Only allows entities from our deduplicated list as nodes  
    entity_set = {e.strip().lower() for e in entity_list}  
    ...
```

```
lines.append(f" {_clean(src)} -- {lbl} --> {_clean(dst)}")
```

- **Why it's clever:** Prevents "garbage nodes" by **strictly enforcing entity validity**.
- **Your takeaway:** Output formats must be *robust* – real data breaks naive assumptions.
- [Mermaid link](#)

Your Assignment (Due in 72 Hours)

Task

Scrape **10 URLs** (provided below), process their text using the DSPy pipeline, and deliver:

- 1. **10 Mermaid diagrams** (one per URL) visualizing key relationships.
- 2. **A structured CSV** with columns: link, tag, tag_type.
- 3. **A Colab notebook** showing your full implementation.

URLs to Scrape

- 1. https://en.wikipedia.org/wiki/Sustainable_agriculture
-

3 / 4

- 2. <https://www.nature.com/articles/d41586-025-03353-5>
- 3. <https://www.sciencedirect.com/science/article/pii/S1043661820315152>
- 4. <https://www.ncbi.nlm.nih.gov/pmc/articles/PMC10457221/>
- 5. <https://www.fao.org/3/y4671e/y4671e06.htm>
- 6. <https://www.medscape.com/viewarticle/time-reconsider-tramadol-chronic-pain-2025a1000ria>
- 7. <https://www.sciencedirect.com/science/article/pii/S0378378220307088>
- 8. <https://www.frontiersin.org/news/2025/09/01/rectangle-telescope-finding-habitable-planets>
- 9. <https://www.medscape.com/viewarticle/second-dose-boosts-shingles-protection-adults-aged-65-years-2025a1000ro7>
- 10. <https://www.theguardian.com/global-development/2025/oct/13/astro-ambassadors-stargazers-himalayas-hanle-ladakh-india>

Assignment Deliverables

1. Mermaid Diagrams (10 total)

- Save as mermaid_{i}.md (e.g., mermaid_1.md)
- **Must include:**
 - o Valid Mermaid syntax (test in [Mermaid Live Editor](#))
 - o Only entities from your deduplicated list as nodes
 - o Edge labels trimmed to 40 chars (as in example)

2. Structured CSV (tags.csv)

link	tag	tag_type
https://...	sustainable agriculture	Concept
https://...	nitrogen uptake	Process
...

Rules:

- tag: Exact entity string (e.g., "pea-barley intercrop", not "intercrop")
- tag_type: Semantic category (e.g., Crop, Process, Measurement)
- **No duplicates per URL** (use your deduplication logic!)

3. Colab Notebook

- **Must include:**
 - o Full code with **comments explaining key steps**
 - o Output csv

4 / 4

How to get Free LLM API key to use with DSPY?

- 1) Follow steps using your own account to get Free LLM api keys from here: https://scribehow.com/viewer/Sign Up for Longcat API Platform_9sYiobPNS0OnXzyxKHu4zg?add_to_team_with_invite=True&sharer_domain=gmail.com&sharer_id=e0b8270f-e494-45b1-b41a-c6adf9f11845
- 2) You might run into limits, so you can ask to increase your limits- https://scribehow.com/viewer/Request More LongCat API Quota_0kdRFILmTdKCulL67qz_rA?add_to_team_with_invite=True&sharer_domain=gmail.com&sharer_id=e0b8270f-e494-45b1-b41a-c6adf9f11845

